



Neural Networks & Pruning

An Implementation from First Principles

Foundations of Machine Learning - Project

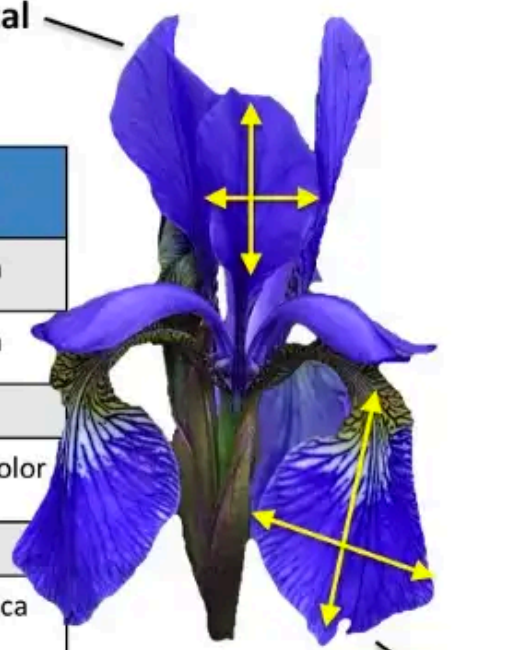
IRIS Dataset

Samples
(instances, observations)

	Sepal length	Sepal width	Petal length	Petal width	Class label
1	5.1	3.5	1.4	0.2	Setosa
2	4.9	3.0	1.4	0.2	Setosa
...					
50	6.4	3.5	4.5	1.2	Versicolor
...					
150	5.9	3.0	5.0	1.8	Virginica

Features
(attributes, measurements, dimensions)

Class labels
(targets)




Number of Samples: 150

Number of Classes: 3

Number of Features: 4

Train - Test Split: 90% - 10%

Digits Dataset Presentation



8 x 8 pixels grayscale images

└─ Number of features 64

Number of Samples: 1797

Number of Classes: 10

Train - Test Split: 90% - 10%

MNIST Dataset Presentation



28 x 28 pixels grayscale images

└→ Number of features 784


Number of Samples: 60 000

Number of Classes: 10

Train - Test Split: 30% - 70%

Implementation Choices

Activation Function



Online Weight Updates:


$$\Delta_{w_{pq}}^{\mu} = \eta \delta_p^{\mu} V_q^{\mu}$$

Loss Function

$$MSE = \frac{1}{2} \sum_{\mu} \sum_i (y_i^{\mu} - O_i^{\mu}(w))^2$$

Training Algorithm

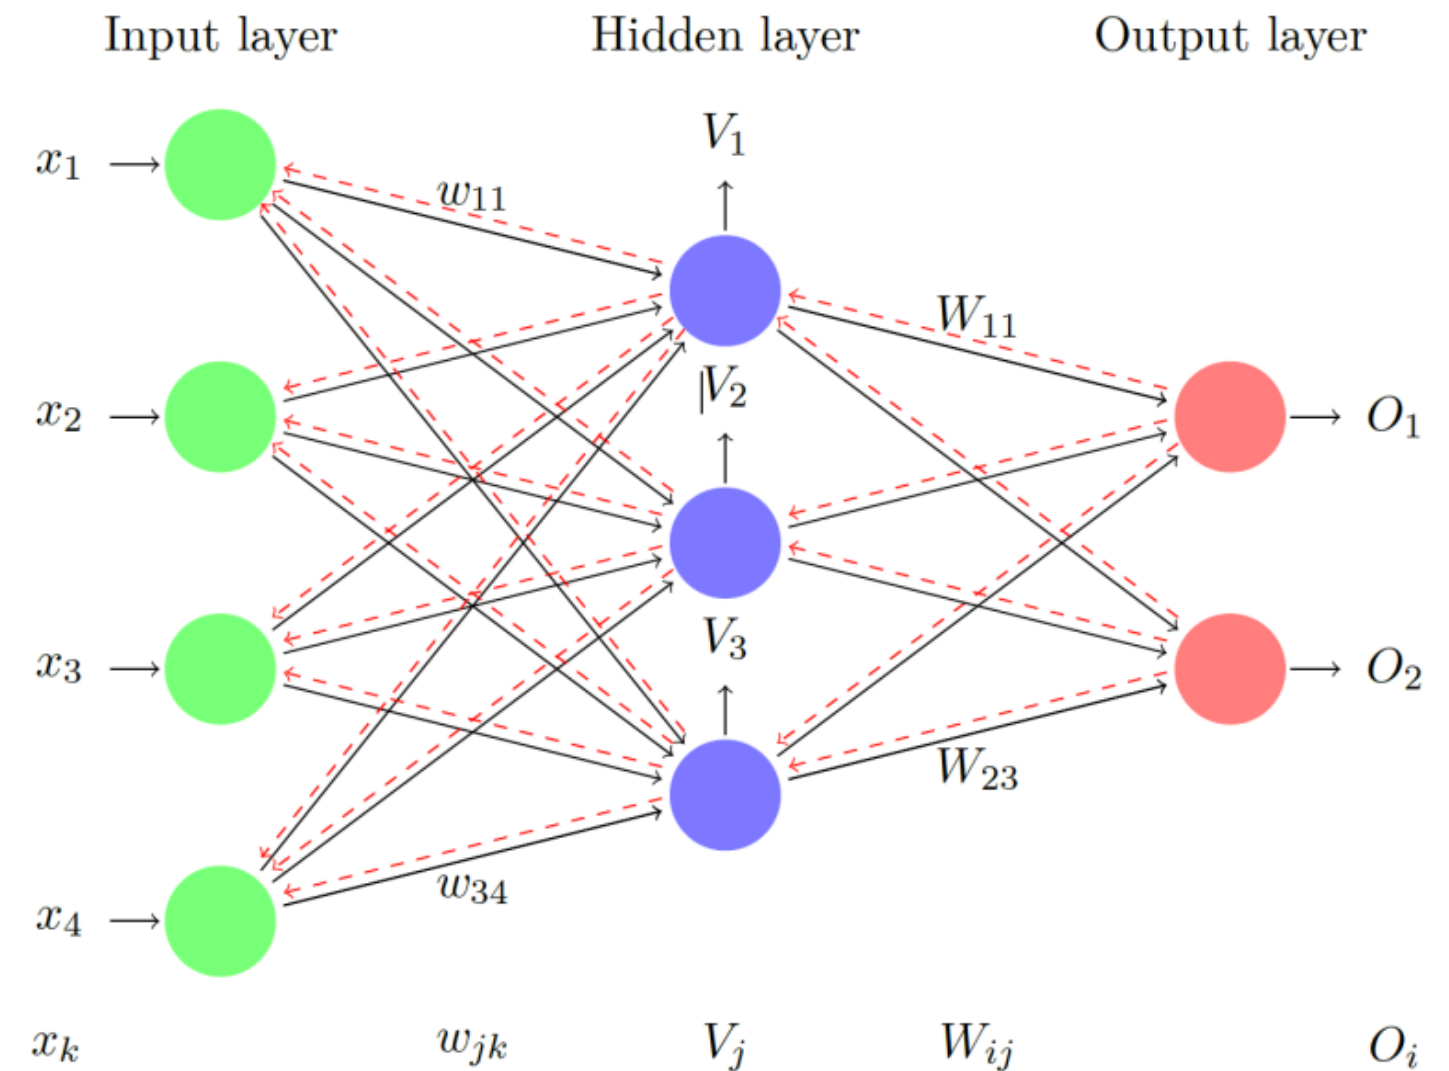
```
def train_procedure(model, epochs, training_set):  
    for epoch in range(epochs):  
        for x, y in range(training_set):  
            forward(model, x)  
            model.compute_deltas(model, y)  
            model.update_weights(model)
```



An example of a Feed-forward Neural Network with one hidden layer (with 3 neurons)

Forward Pass

```
def forward(model, x):  
    for layer in model_layer:  
        extend_x_with_bias(x, -1)  
  
        # compute non activated input  
        h = dot_prod(x, layer)  
        # applicate non linear activation  
        V = layer.activate(h)  
  
        model.h.append(h)  
        model.V.append(V)  
  
    return model.V.last_value()
```



Weights Update

```
def update_weights(model):  
    for i, layer in model.layers:  
        if i == 0: # retrieve the preceding activated output  
            V = model.input # in the first layer we use input  
        else:  
            V = model.V[i-1]  
  
        layer += model.deltas[i] * V # update
```


Delta Computation

```
def compute_deltas(model, y):
    # traverse the network in reverse
    for i, layer in reversed(model.layers):
        if hidden_to_output(layer):
            d_ij = layer.activate_derivative(model.h[i])
            model.delta.append((y-model.V[i])) * d_ij
        else:
            d_jk = layer.activate_derivative(model.h[i])



            previous_w = model.layers[i+1]
            previous_delta = model.delta[i+1]

            previous_weighted_error = dot_product(previous_w, previous_delta)

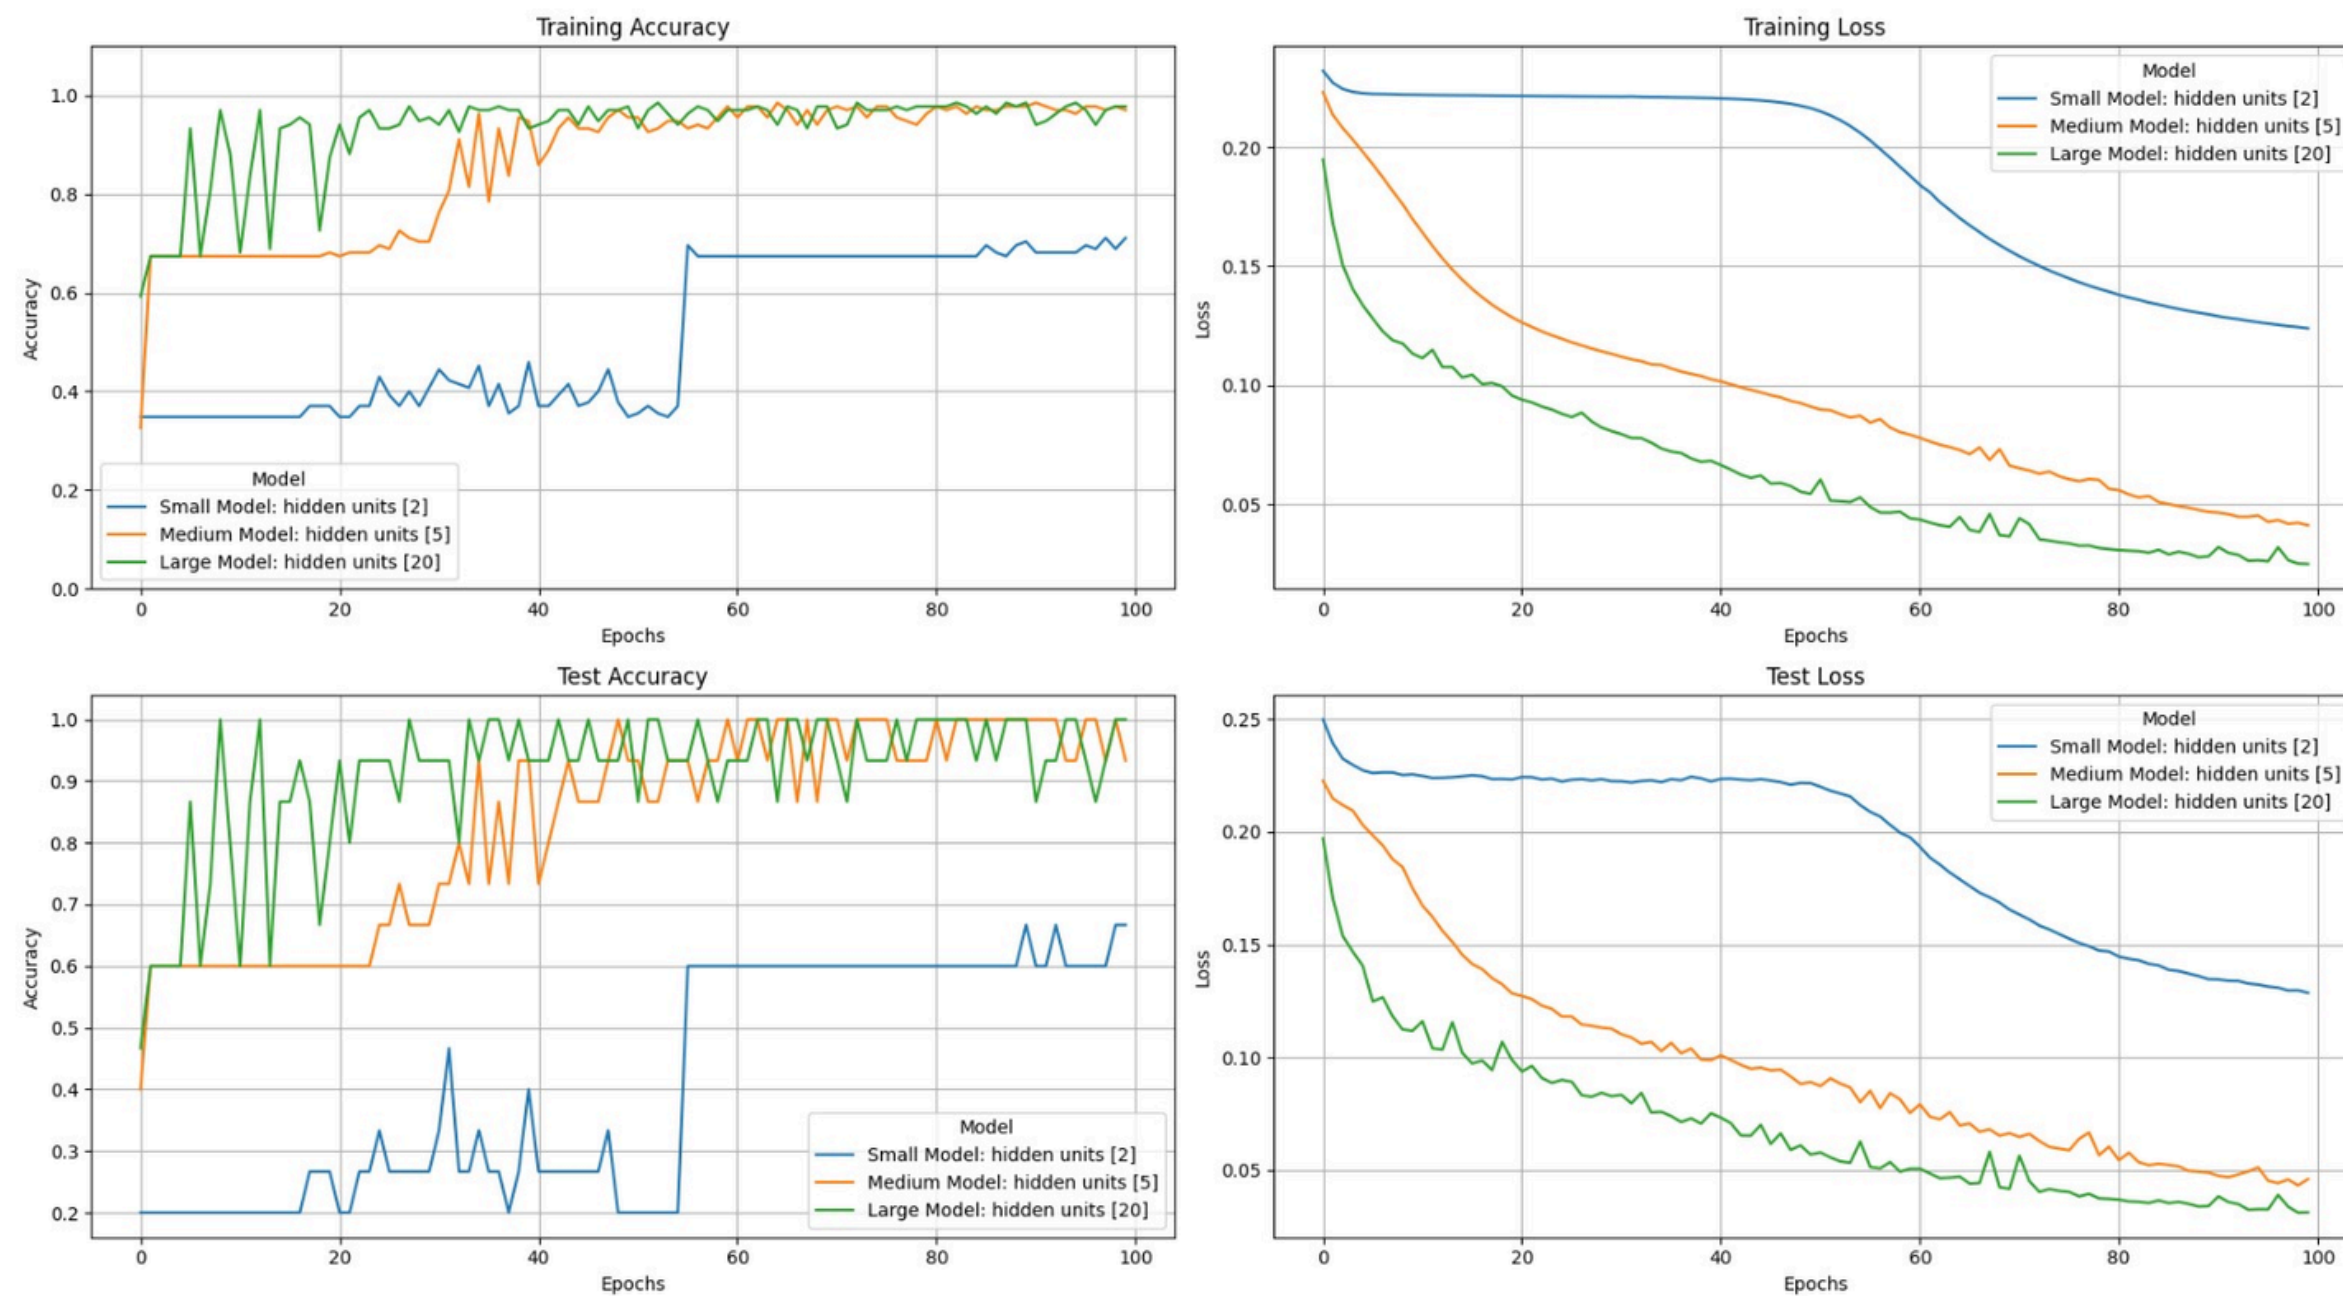
            d_jk *= previous_weighted_error
            model.delta.append(d_jk)
```

$$\delta_i^\mu = (\text{target}_i^\mu - y_i^\mu) g'(h_i^\mu)$$

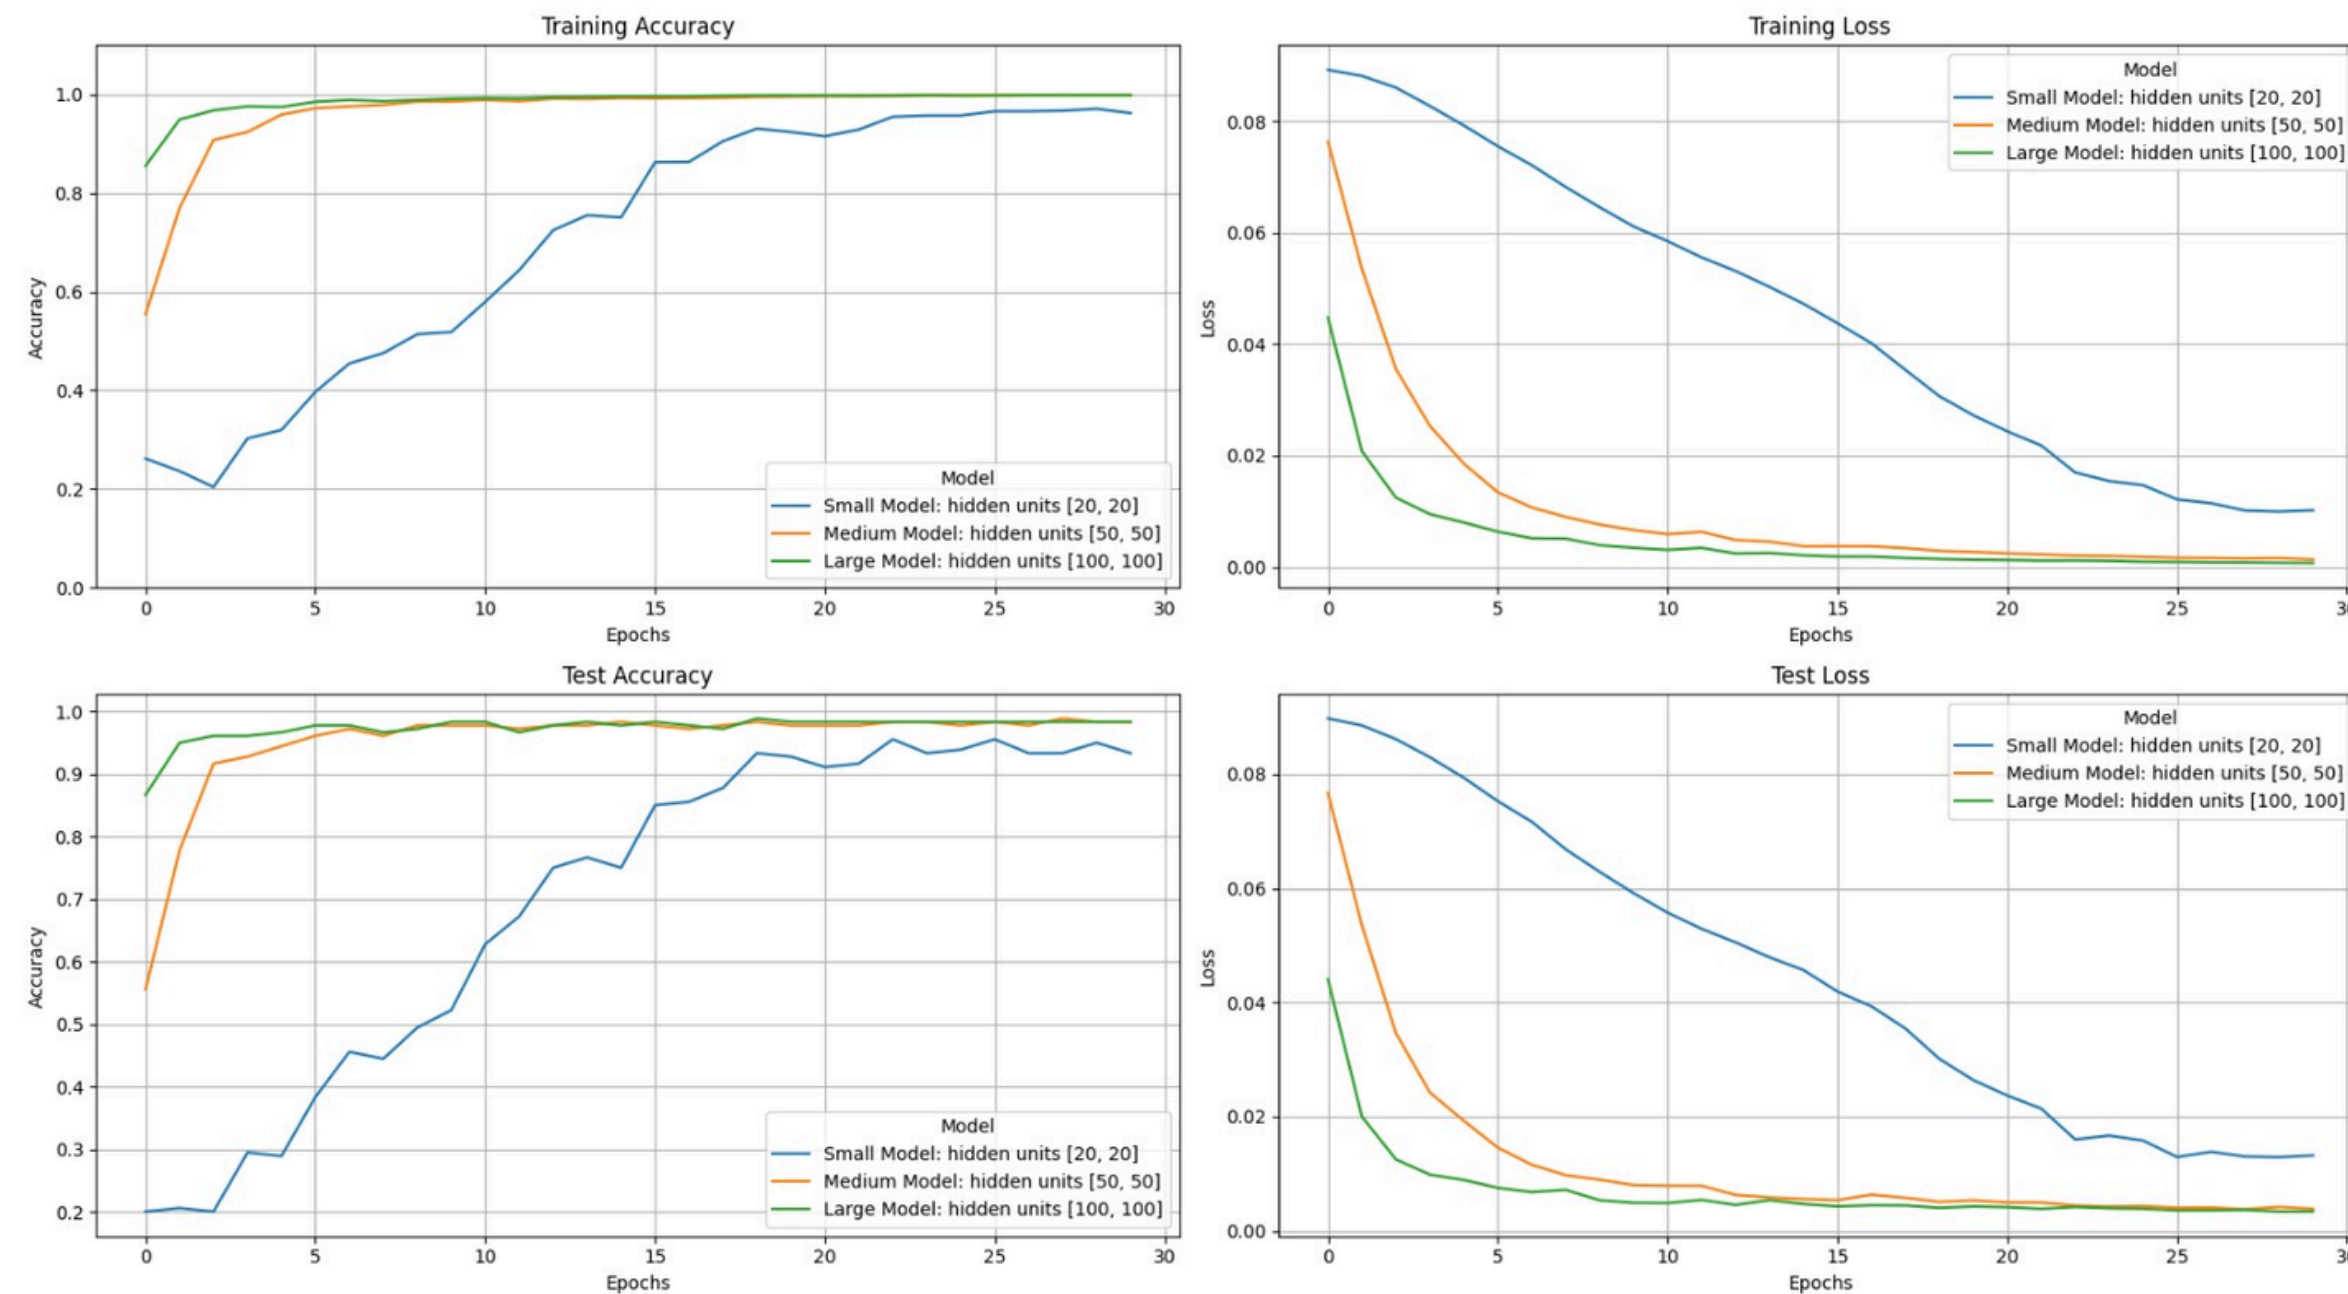
$$\hat{\delta}_j^\mu = g'(h_j^\mu) \sum_i \delta_i^\mu W_{ij}$$




Performance on IRIS Dataset



Performance on DIGITS Dataset




Performance on MNIST Dataset



Iterative Pruning Algorithm

Hidden unit removal
selection

Update of the remaining
hidden unit connections to
to output layer




Pruning Algorithm - Removal Selection

```
def choose_unit_to_prune(model):  
    # shape: (n_train_samples * n_hidden_units)  
    hidden_layer_output = model.batch_hidden_layer_output()  
    hidden_to_output_weights = model.weights[-1]  
  
    norms = [0 for _ in range(model.n_hidden_units)]  
    for h in range(model.n_hidden_units):  
        # final length = n_output_units * n_input_samples  
        output_units_b = []  
  
        for output_idx in range(model.n_outputs):  
            # one b_i for each train sample  
            b_i = hidden_to_output_weights[h,i] * hidden_layer_out[:,h]  
            output_units_b.append(b_i)  
  
        norms.append(compute_norm(output_units_b))  
  
    return argmin(norms)
```

We compute the **contribution** of each **hidden unit** to each **output unit** for every input sample

The contribution is given by the connection weight and the h-th hidden unit's output



Pruning Algorithm - Update Existing Connections

```
def update_connections(model, h_star):
    hidden_output_without_removed = model.hidden_output_removal(h_star)
    deltas = zeros_with_shape(model.n_output_units, model.n_hidden_units-1)

    for i in range(model.output_unit_count):
        # contribution of the removed hidden to the output unit i
        b_i = model.w[-1][i, h_star] * model.hidden_output[:, h_star]

        delta_i = solve_with_least_squares(hidden_output_without_removed, b_i)


        deltas[i] = delta_i

    model.update_hidden_to_output_weights(deltas)
    model.remove_hidden_layer_input(h_star)
```

We solve an **over-determined** system of equations for each output unit touched by the weight update

$$\underbrace{\begin{bmatrix} y_{11} & y_{12} & \cdots & a_{1,n_{\text{hidden}}} \\ y_{21} & y_{22} & \cdots & y_{2,n_{\text{hidden}}} \\ \vdots & \vdots & \ddots & \vdots \\ y_{n_{\text{samples}},1} & a_{n_{\text{samples}},2} & \cdots & a_{n_{\text{samples}},n_{\text{hidden}}} \end{bmatrix}}_{Y \in \mathbb{R}^{n_{\text{samples}} \times n_{\text{hidden}}}} \underbrace{\begin{bmatrix} x_1 \\ x_2 \\ \vdots \\ x_{n_{\text{hidden}}} \end{bmatrix}}_{x \in \mathbb{R}^{n_{\text{hidden}}}} = \underbrace{\begin{bmatrix} b_1 \\ b_2 \\ \vdots \\ b_{n_{\text{samples}}} \end{bmatrix}}_{b \in \mathbb{R}^{n_{\text{samples}}}}$$

Performance on DIGITS Dataset





Neural Networks & Pruning

An Implementation from First Principles

Foundations of Machine Learning - Project